

FALL 2026

AI SAFETY HACKATHON

What is the model actually doing?

Open the model. Stress-test the system. Build the evaluation.
Bring a question or start from a scaffolded problem.

INTERPRETABILITY

SAFETY EVALUATION

CYBERSECURITY

KICKOFF Friday, Oct 9

9:00–11:00 AM
Capital One Hub
Theme, tools, and team formation

BUILD DAY + TALKS Friday, Oct 16

9:00 AM–5:00 PM
Capital One Hub

FEATURED KEYNOTE SPEAKER

Greg Frank

Co-founder and Chief
Scientist
MoltAI
Oct 16 · 9:30 AM

REGISTER BY SEPTEMBER 25

**No team or prior interpretability
experience required.**

Teams of 2–4; solo sign-ups matched. Starter code, GPU
credits, and five-minute team presentations.

PRESENTED BY



DETAILS + REGISTRATION

